Stroke Prediction Dataset

Classification Project

Tomer Badug. Shirli Miller. Judi Eliya

project goal

Prediction Stroke

what is a stroke?

Stroke is a medical emergency. A stroke occurs when blood flow to a part of your brain is interrupted or reduced, preventing brain tissue from getting oxygen and nutrients. Brain cells begin to die within minutes

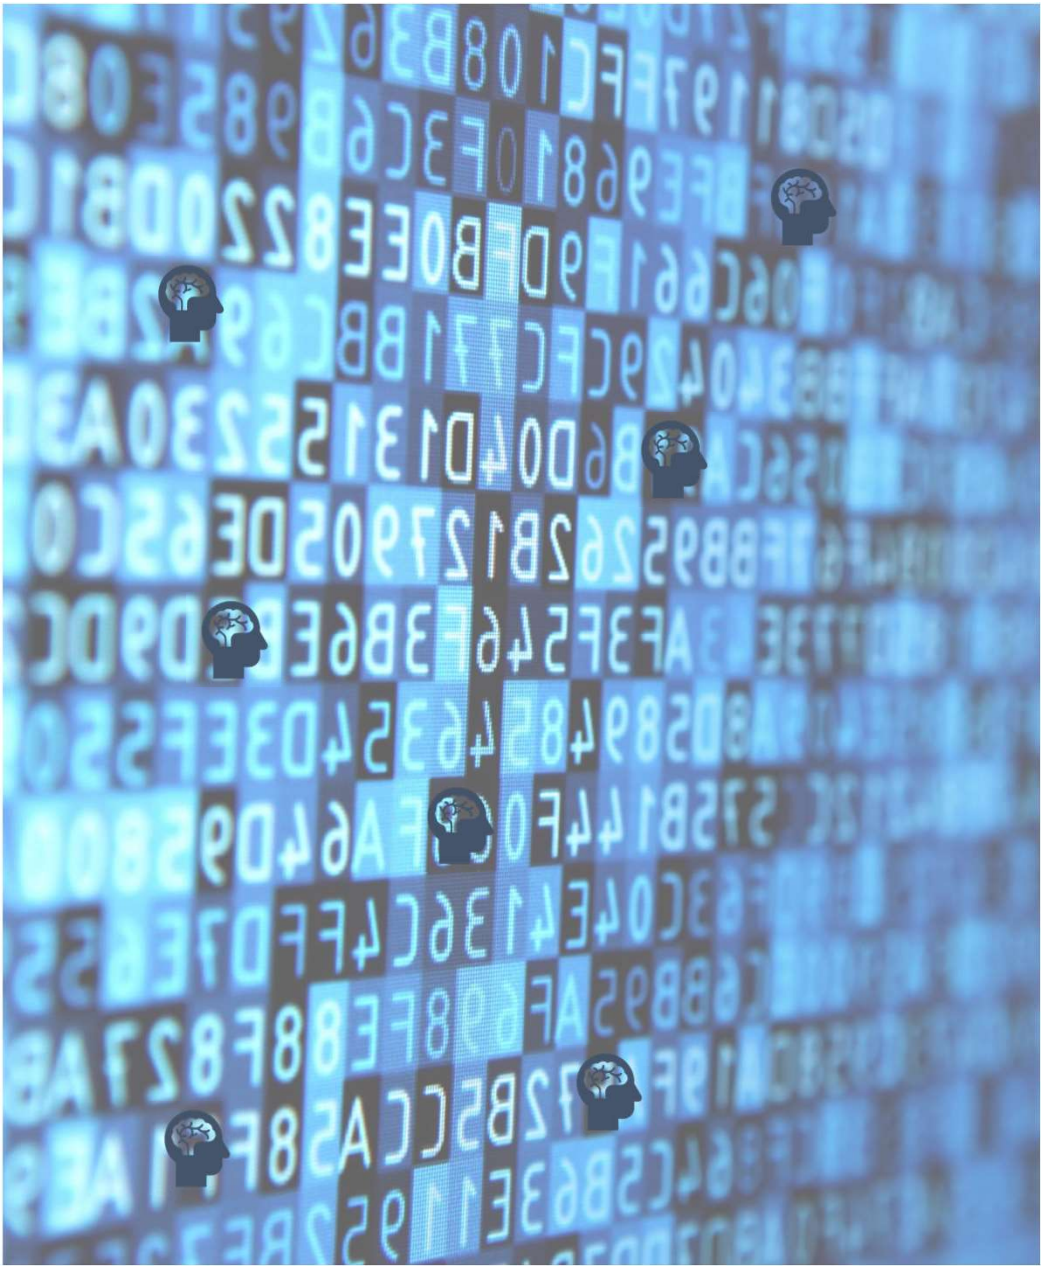

Data content

Feature:

1. id
2. gender
3. age
4. hypertension
5. heart_disease
6. ever_married
7. work_type
8. Residence_type
9. avg_glucose_level
10. bmi
11. smoking_status

12 columns
5110 row

Target variable – stroke

Exploratory data analysis -EDA

- Missing Values Handling
- Remove duplicate and outlier
- Feature selection
- Compute pairwise correlation of columns – matrix
- Encoding
- Understanding the variables

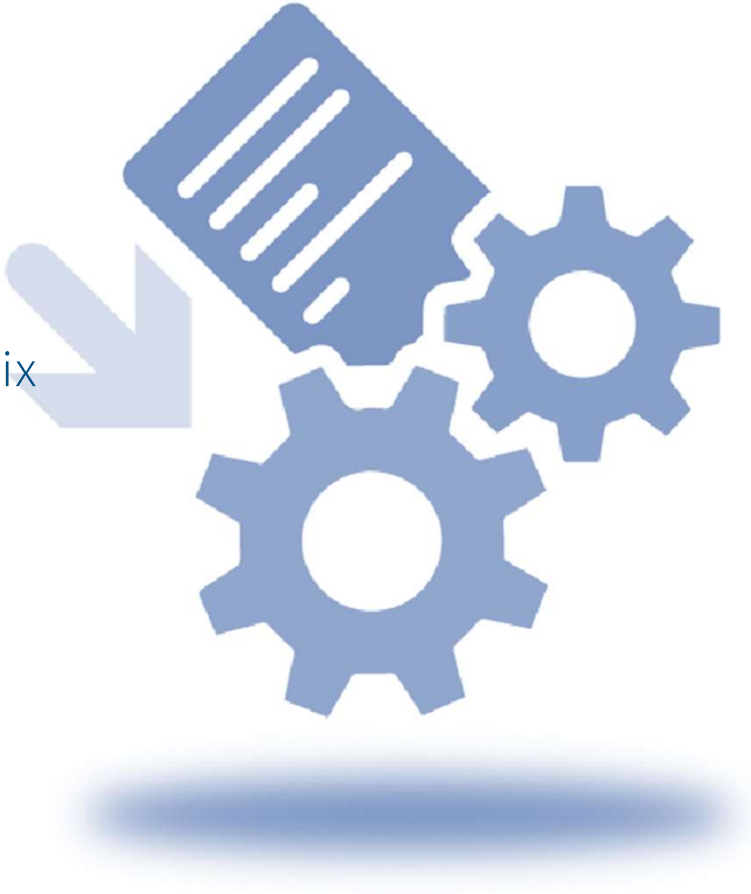

EDA

1. ID

Feature selection

1. Unique value
2. Drop

11 columns, 5110 row

EDA

2. Gender

"Male", "Female" or "Other"

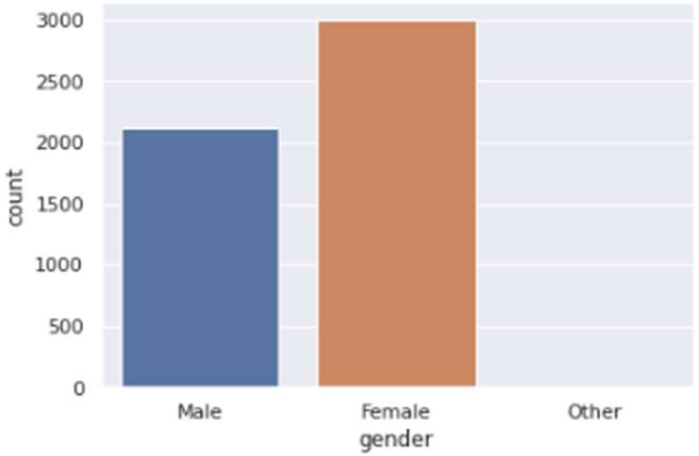

Female	2994
Male	2115
Other	1

	srtoke	0	1	
gender	**Female**	2853	141	4.71%
	Male	2007	108	5.11%
	Other	1	0	0.00%

11 columns, 5109 row

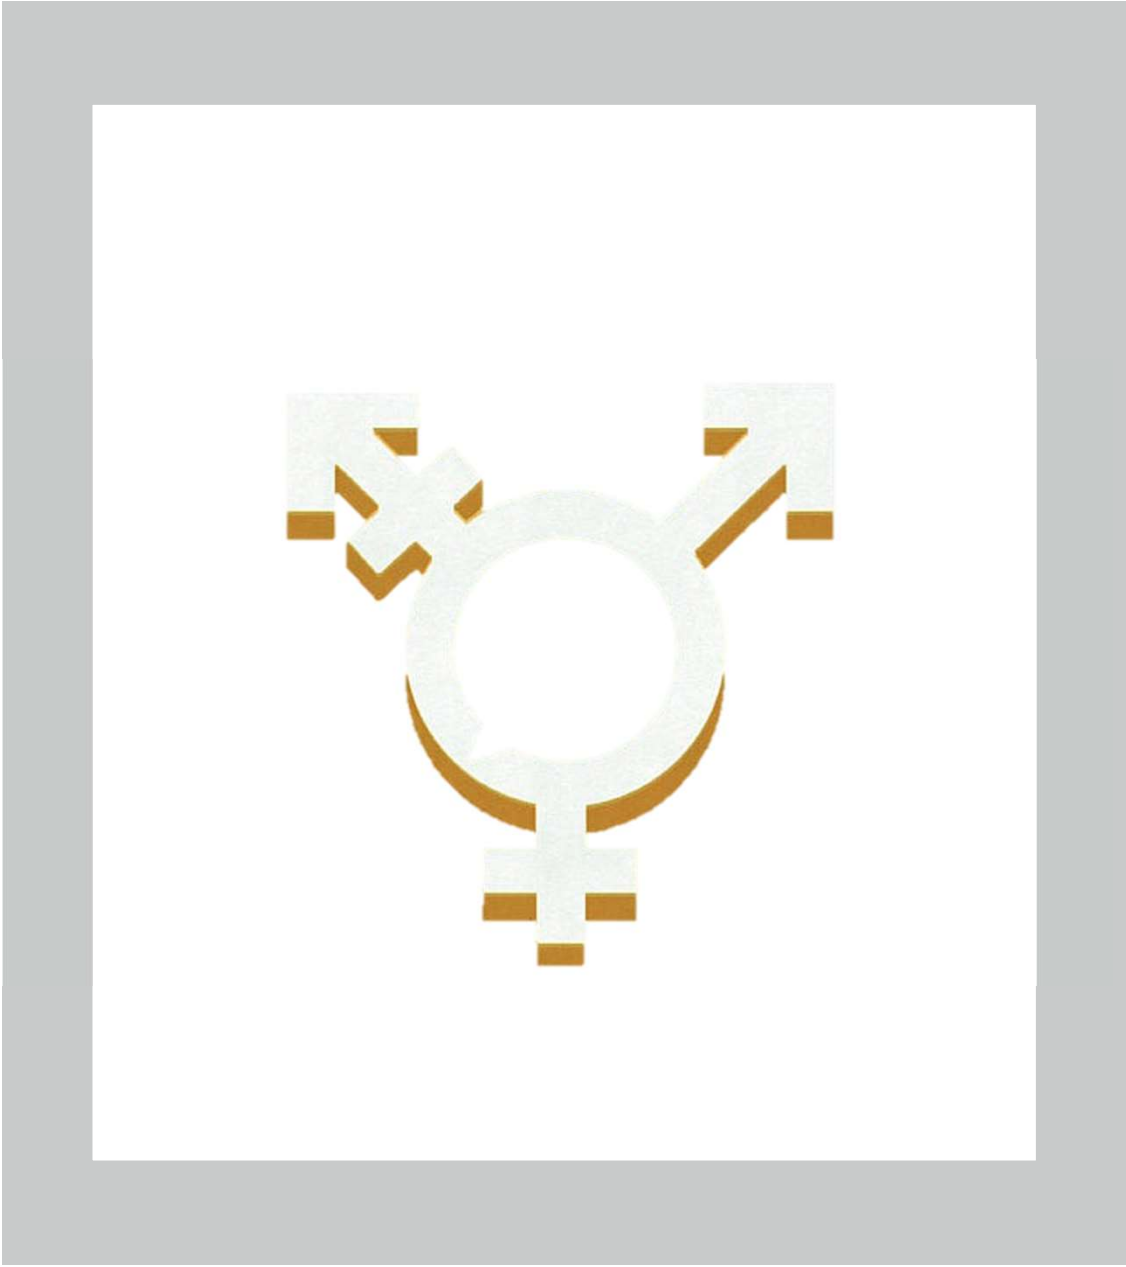

EDA

3. Age:

Age of the patient

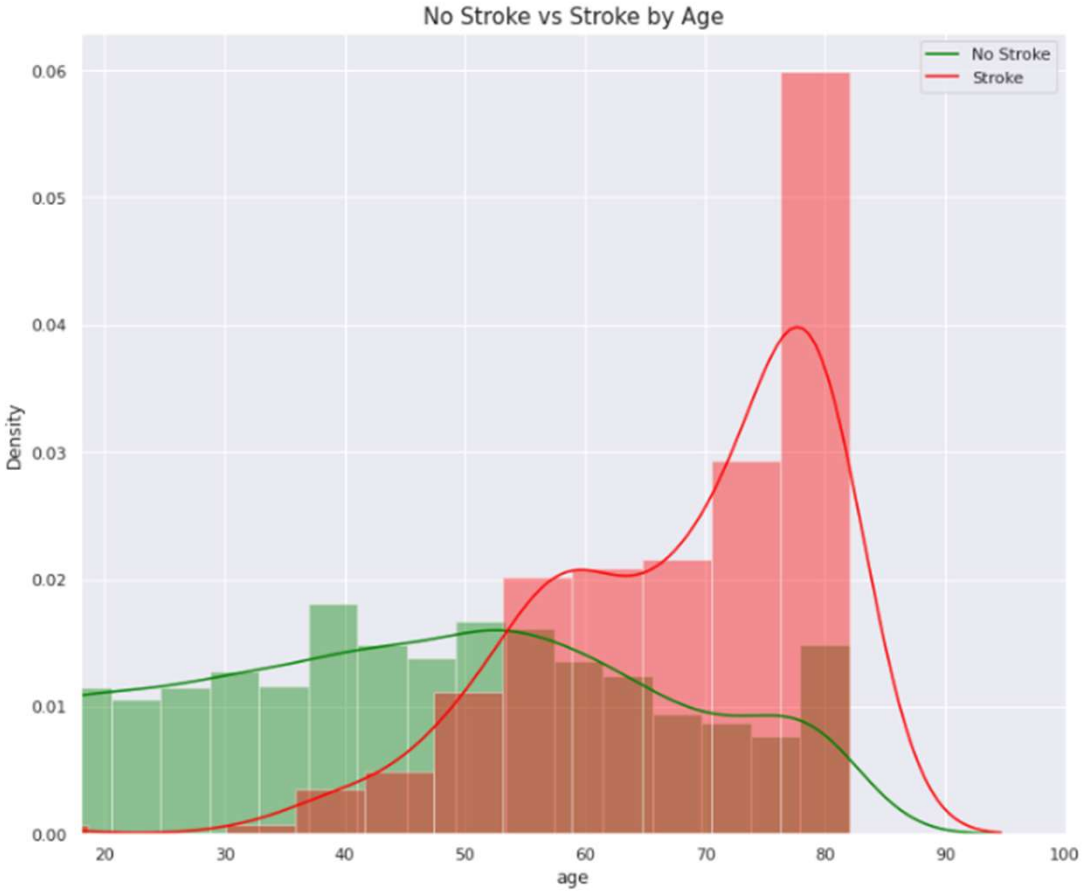

11 columns, 5109 row

	age
count	5110.000000
mean	43.226614
std	22.612647
min	0.080000
25%	25.000000
50%	45.000000
75%	61.000000
max	82.000000

EDA

4. Hypertension :

0 - No hypertension

1 - Hypertension

	srtoke	0	1	
hypertension	**0**	4429	183	3.97%
	1	432	66	13.25%

11 columns, 5109 row

EDA

5. Heart disease :

0 - No heart diseases

1 - Heart disease

		srtoke	0	1
heart_disease	0	4632	202	4.18%
	1	229	47	17.03%

11 columns, 5109 row

EDA

6. Ever married :

"No" or "Yes"

	srtoke	0	1	
ever_married	**No**	1728	29	1.65%
	Yes	3133	220	6.56%

11 columns, 5109 row

EDA

7. Work type :

"Private"

"Self-employed“

"Govt_jov“

"Children“

"Never_worked"

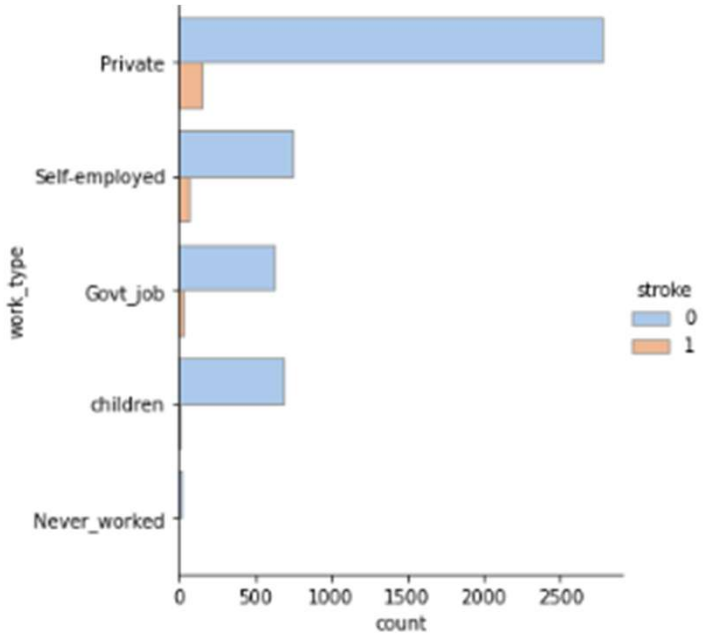

	srtoke	0	1	
work_type	Govt_job	624	33	5.02%
	Never_worked	22	0	0.00%
	Private	2776	149	5.09%
	Self-employed	754	65	7.94%
	children	685	2	0.29%

11 columns, 5109 row

EDA

9. Residence type :

"Rural" or "Urban"

	srtoke	0	1	
Residence_type	**Rural**	2400	114	4.53%
	Urban	2461	135	5.20%

11 columns, 5109 row

EDA

9. Avg Glucose level:

Average glucose level in blood

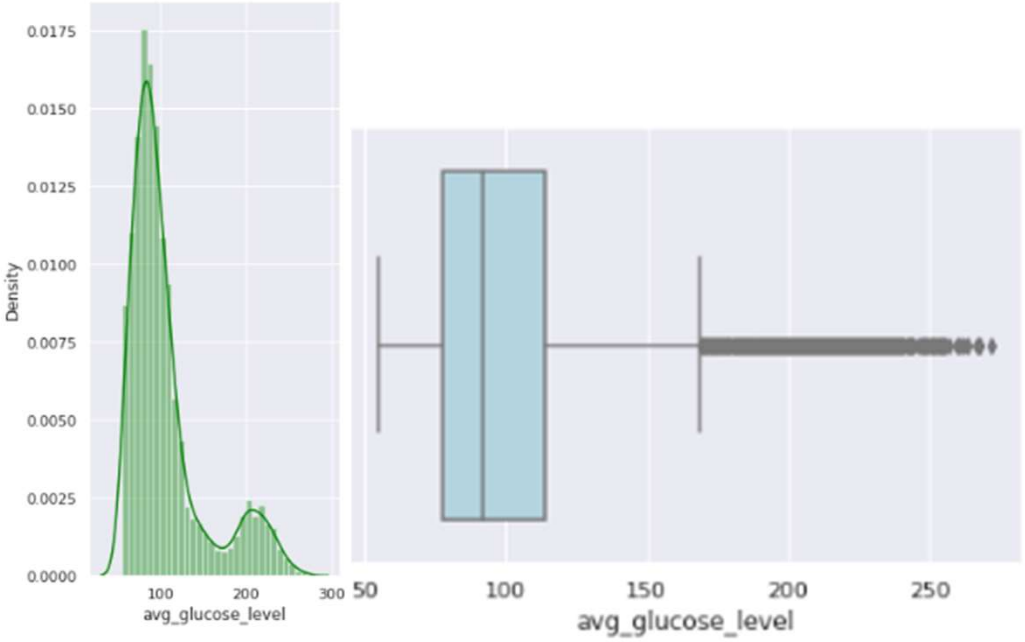

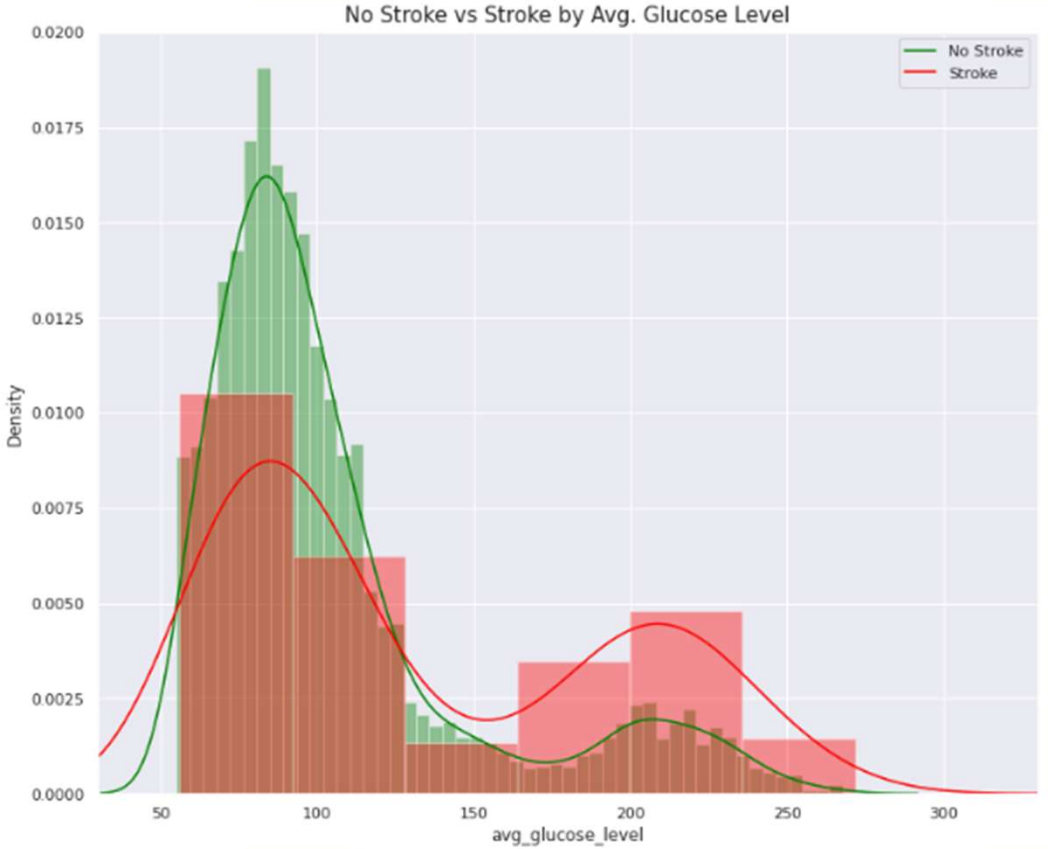

11 columns, 5109 row

EDA

10. BMI:

Body Mass Index

201 null - Replacing the missing values with mean

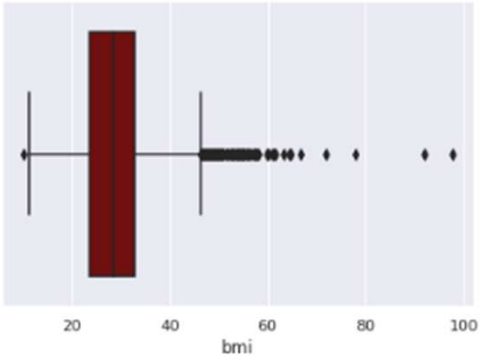

Outlier:

Decision - Drop 5 max bmi

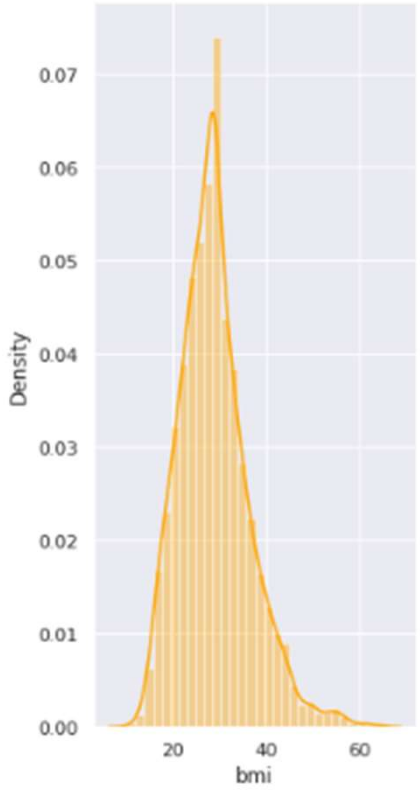

11 columns, 5104 row

EDA

10. Smoking Status:

"Formerly smoked“
"Never smoked“
"Smokes“
"Unknown"

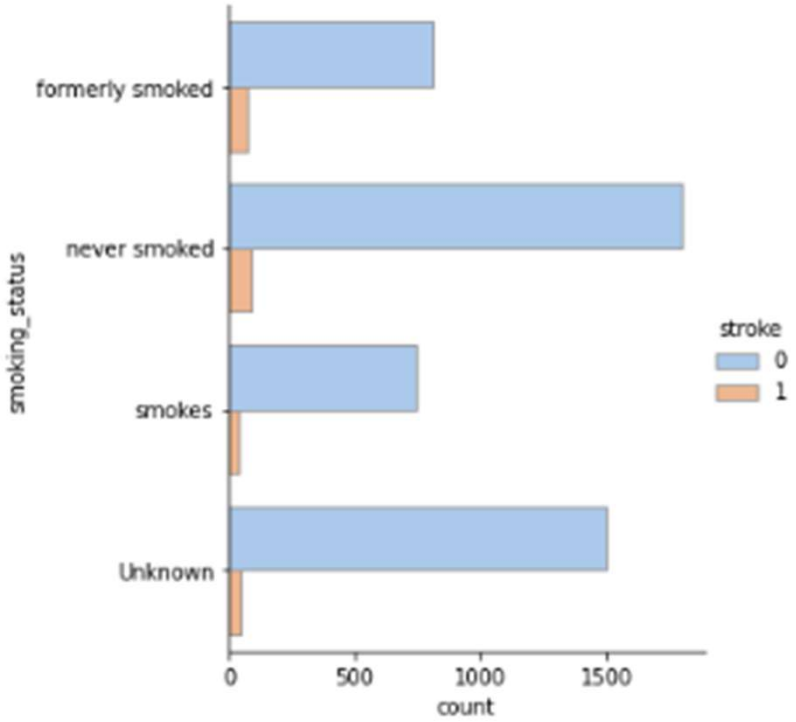

	srtoke	0	1	
smoking_status	**Unknown**	1497	47	3.04%
	formerly smoked	815	70	7.91%
	never smoked	1802	90	4.76%
	smokes	747	42	5.32%

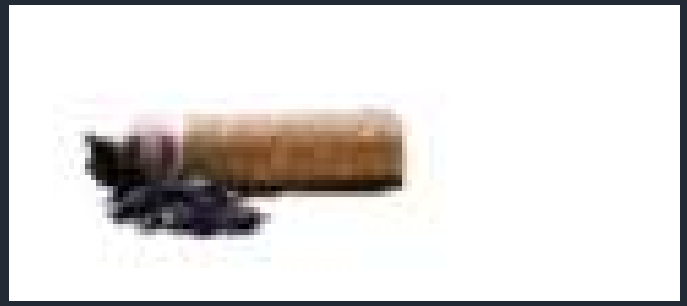

Target variable

12. Stroke:

0 = No stroke

1 = Stroke

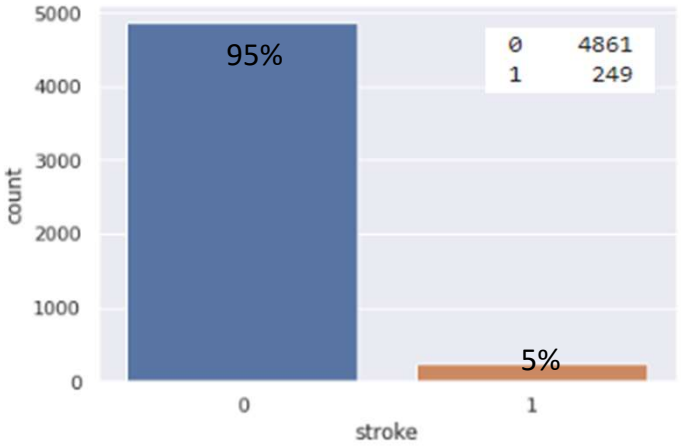

Conclusion and the next steps

Age has the highest impact on stroke, even though the stroke also depends on the other variables,

Such as: glucose level, heart disease, blood pressure, smoking and even type of work and area of residence

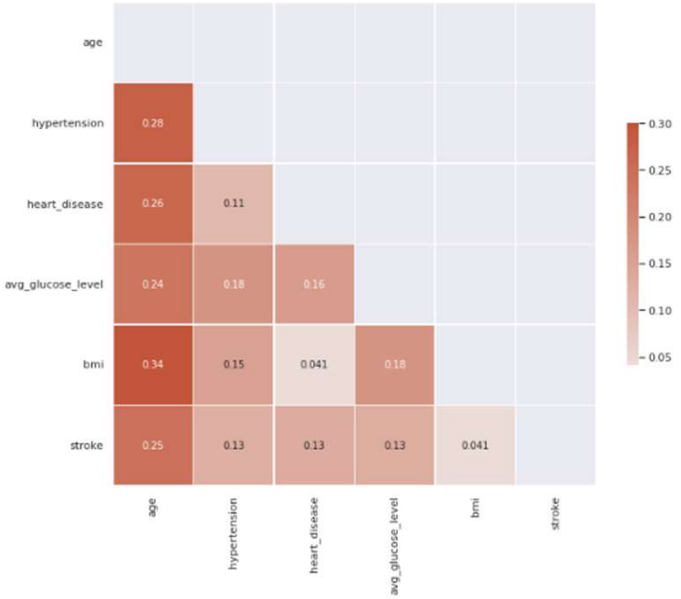

Conclusion and the next steps

Standard Scaler → Splitting the dataset and Handling imbalanced → Create classification models → Valuation of results → Comparison → Future research

Standard Scaler

- The models With and Without Standart Scaler

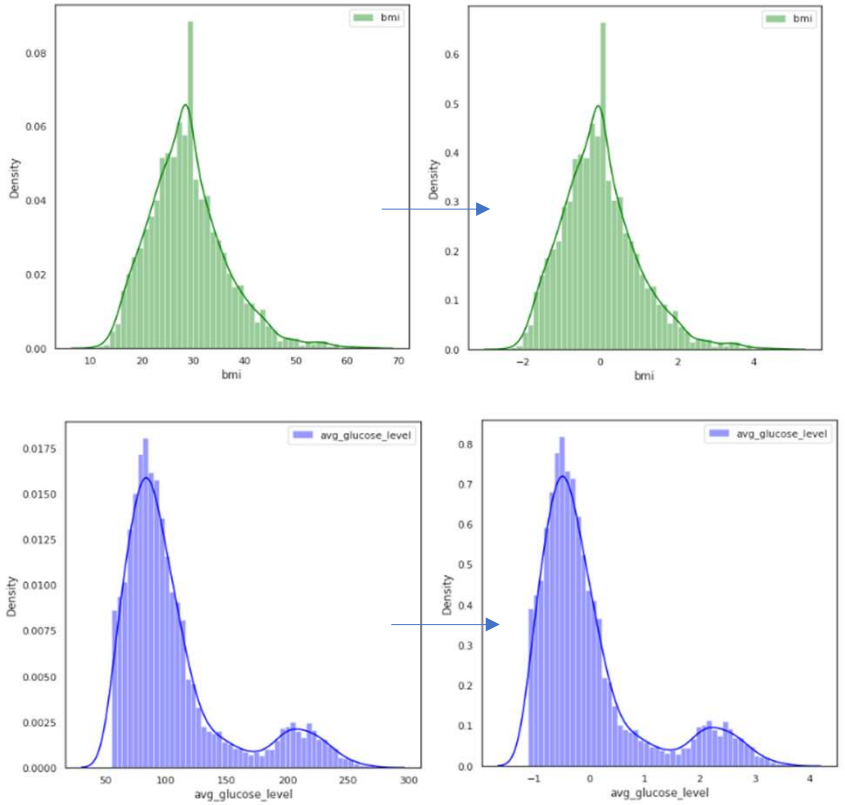

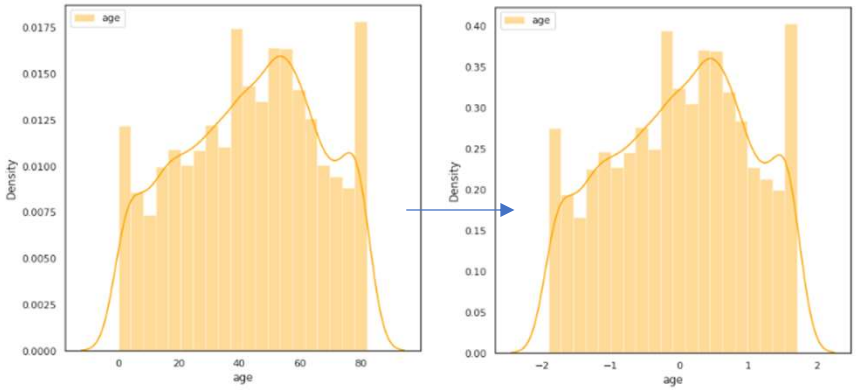

Splitting data and Handling imbalanced

- Splitting dataset 30-70
- stratify (95-5)
- Over Sampling using SMOTE (50-50 train data)

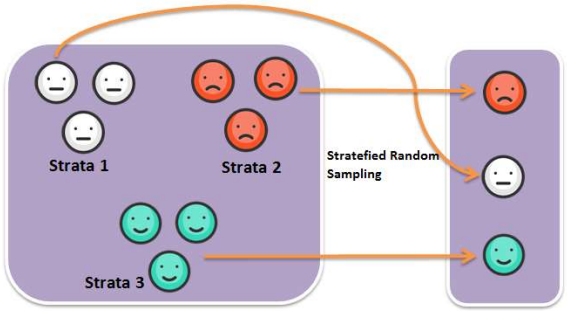


```
Before OverSampling, counts of label '1': 174
Before OverSampling, counts of label '0': 3398

After OverSampling, the shape of train_X: (6796, 17)
After OverSampling, the shape of train_y: (6796,)

After OverSampling, counts of label '1': 3398
After OverSampling, counts of label '0': 3398
```

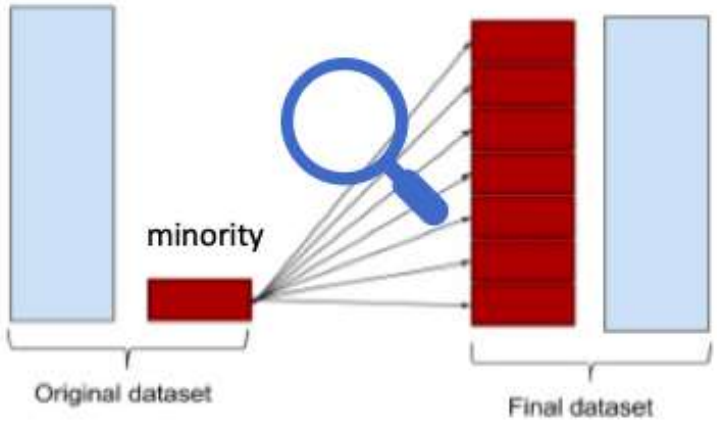

classification models

- Logistic Regression
- SVM
- Random Forest

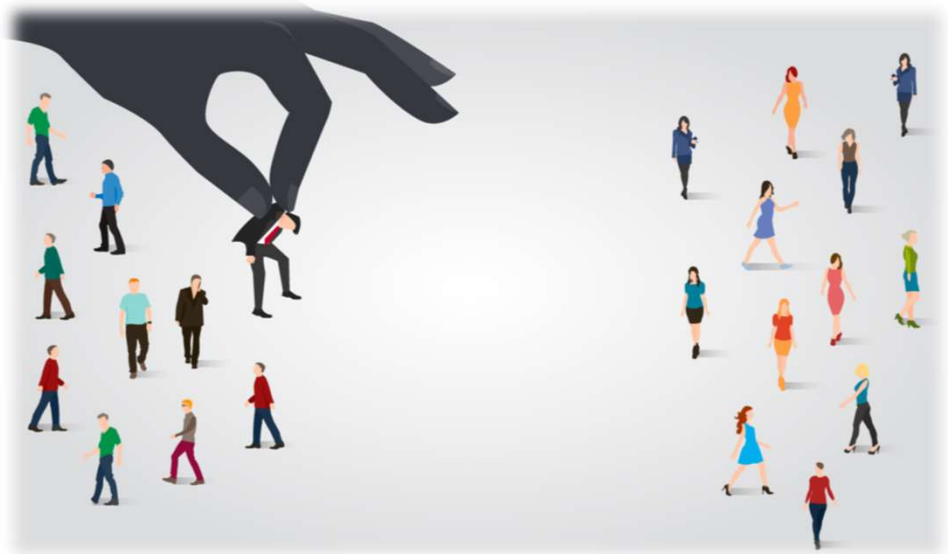

valuation of results

- Precision
- Recall
- F1

and connection

Logistic Regression

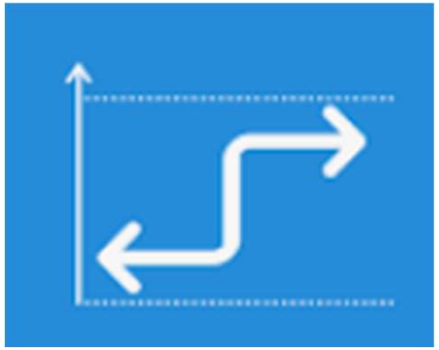

- With and without Over Sampling

Testing Score
0.9510443864229765

without OS

	precision	recall	f1-score	support
0	0.95	1.00	0.97	1457
1	0.00	0.00	0.00	75
accuracy			0.95	1532
macro avg	0.48	0.50	0.49	1532
weighted avg	0.90	0.95	0.93	1532

	0	1
0	1457	0
1	75	0

Testing Score
0.7271540469973891

with OS

	precision	recall	f1-score	support
0	0.99	0.72	0.83	1457
1	0.13	0.79	0.22	75
accuracy			0.73	1532
macro avg	0.56	0.76	0.53	1532
weighted avg	0.94	0.73	0.80	1532

	0	1
0	1055	402
1	16	59

SVM

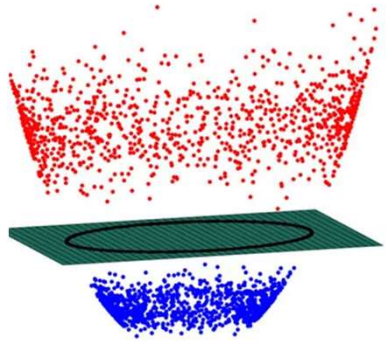

- With and without Over Sampling
- With and without Standart Scaler

- Cross Validation

```
Scores : 0.802 0.781 0.790 0.776 0.797 0.780 0.795
```

- Grid Search

```
Best parameters set found on development set:

{'C': 1000, 'gamma': 0.001, 'kernel': 'rbf'}
```

	precision	recall	f1-score	support
0	0.98	0.76	0.85	1457
1	0.13	0.68	0.21	75
accuracy			0.75	1532
macro avg	0.55	0.72	0.53	1532
weighted avg	0.94	0.75	0.82	1532

```
Training Score 0.796203649205415
Testing Score 0.7088772845953003
```

	0	1
0	1026	431
1	15	60

Random Forest

- With and without Over Sampling
- Grid Search

```
Best Parameters :  {'criterion': 'gini', 'n_estimators': 150, 'random_state': 0}
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	1457
1	0.09	0.01	0.02	75
accuracy			0.95	1532
macro avg	0.52	0.50	0.50	1532
weighted avg	0.91	0.95	0.93	1532

	0	1
0	1447	10
1	74	1

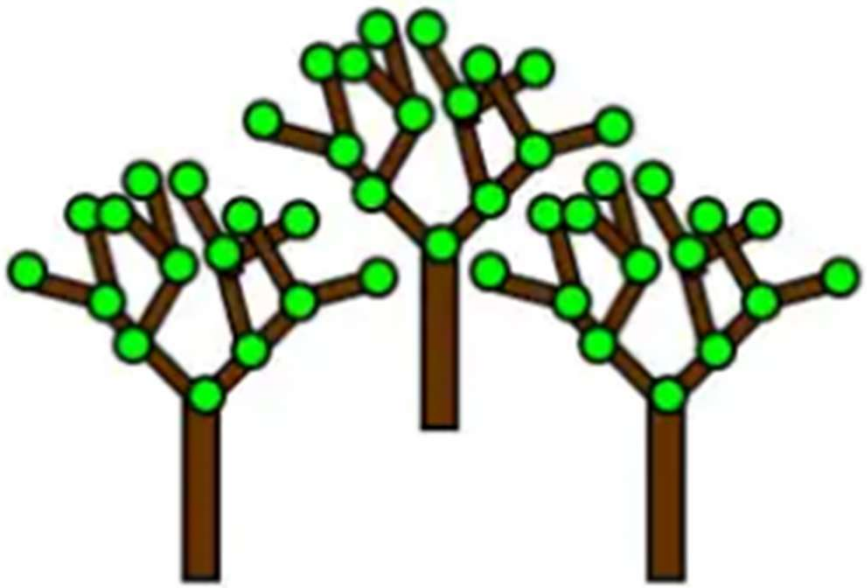

Present main results and comparison of methods

Grid Search

scoring = 'accuracy'

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                   intercept_scaling=1, l1_ratio=None, max_iter=100,
                   multi_class='auto', n_jobs=None, penalty='l2',
                   random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                   warm_start=False):
Best Accuracy : 77.27%
Best Parameters :  {'C': 0.5, 'random_state': 0}

----------------

SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False):
Best Accuracy : 78.13%
Best Parameters :  {'C': 0.5, 'kernel': 'linear', 'random_state': 0}

----------------

RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                       criterion='gini', max_depth=None, max_features='auto',
                       max_leaf_nodes=None, max_samples=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=1, min_samples_split=2,
                       min_weight_fraction_leaf=0.0, n_estimators=100,
                       n_jobs=None, oob_score=False, random_state=None,
                       verbose=0, warm_start=False):
Best Accuracy : 97.29%
Best Parameters :  {'criterion': 'gini', 'n_estimators': 100, 'random_state': 0}
```

scoring = 'recall'

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                   intercept_scaling=1, l1_ratio=None, max_iter=100,
                   multi_class='auto', n_jobs=None, penalty='l2',
                   random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                   warm_start=False):
Best Accuracy : 81.64%
Best Parameters :  {'C': 0.5, 'random_state': 0}

----------------

SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False):
Best Accuracy : 85.17%
Best Parameters :  {'C': 1, 'kernel': 'linear', 'random_state': 0}

----------------

RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                       criterion='gini', max_depth=None, max_features='auto',
                       max_leaf_nodes=None, max_samples=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=1, min_samples_split=2,
                       min_weight_fraction_leaf=0.0, n_estimators=100,
                       n_jobs=None, oob_score=False, random_state=None,
                       verbose=0, warm_start=False):
Best Accuracy : 95.29%
Best Parameters :  {'criterion': 'gini', 'n_estimators': 150, 'random_state': 0}
```

Present main results and comparison of methods

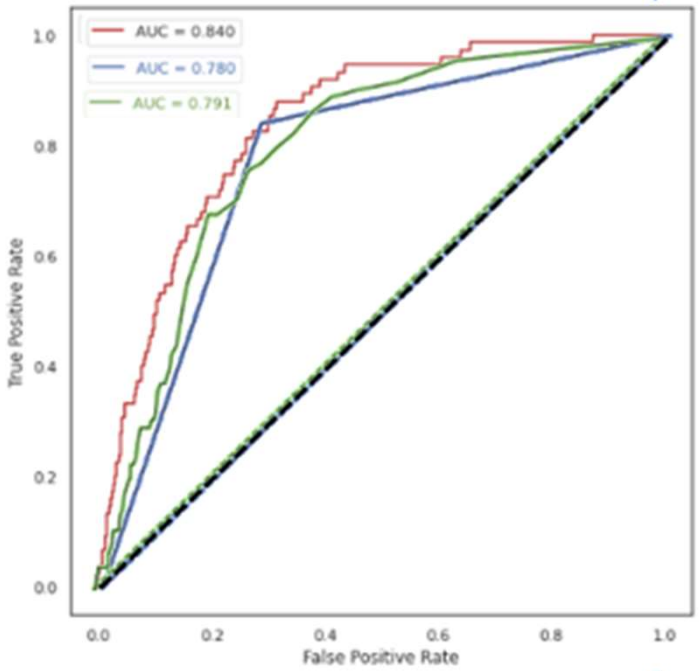

scoring = 'recall'

Logistic Regression

	P	N
P	1073	384
N	14	61

	precision	recall	f1-score	support
0	0.99	0.74	0.84	1457
1	0.14	0.81	0.23	75
accuracy			0.74	1532
macro avg	0.56	0.77	0.54	1532
weighted avg	0.95	0.74	0.81	1532

SVM

	P	N
P	1050	407
N	12	63

	precision	recall	f1-score	support
0	0.99	0.72	0.83	1457
1	0.13	0.84	0.23	75
accuracy			0.73	1532
macro avg	0.56	0.78	0.53	1532
weighted avg	0.95	0.73	0.80	1532

random forest

	P	N
P	1448	9
N	74	1

	precision	recall	f1-score	support
0	0.95	0.99	0.97	1457
1	0.10	0.01	0.02	75
accuracy			0.95	1532
macro avg	0.53	0.50	0.50	1532
weighted avg	0.91	0.95	0.93	1532

summary and Future research

- Scaling VS Threshold
- Defining additional / other variables in Grid Search
- Examination of additional models
- Combine scores and determining weight for each of them
- ensemble methods

Thanks

Tomer badug

Shirli miller

Judi Eliya